

1. Answer the questions in the collaboration policy on page 1.

- Did you receive any help whatsoever from anyone in solving this assignment?

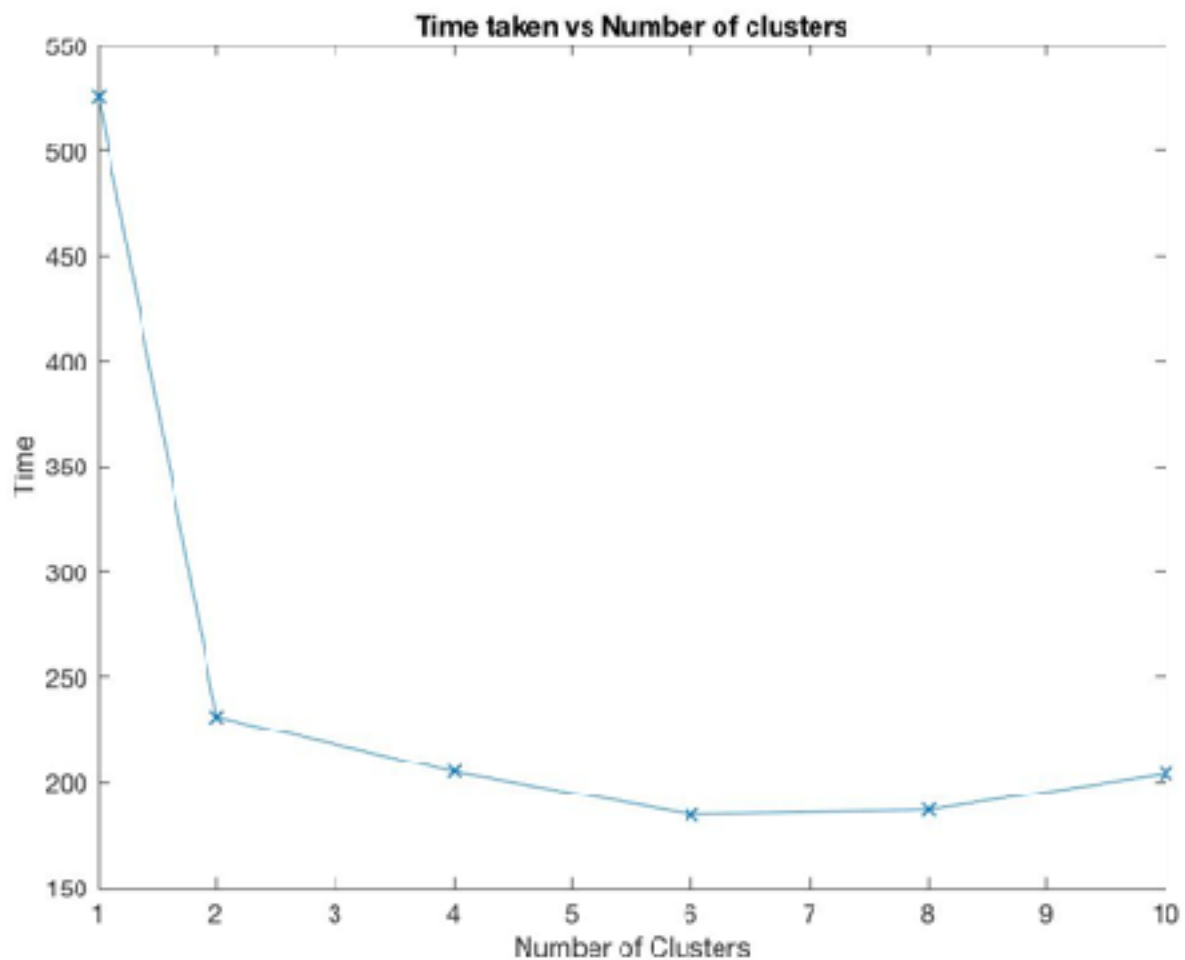
NO

- Did you give any help whatsoever to anyone in solving this assignment?

NO

2. As we increase number of clusters, time taken to complete the reduce process decreases upto some limit (6 clusters), after that it starts increasing. There can be many reasons behind this:

- A. Communication overhead : The output of sort operation has to be sent to reducers, as number of reducers increase this time increases. If data set is large enough this time won't have much effect on total time because of computation overhead. But after certain limit, it will overweight and become more than computation time.
- B. If every processor in reduce step, does not have same speed or amount of work to do, (data is not evenly distributed across processes), the time taken for whole operation will be determined by processor that has smallest data.



3.

Mapper 1:

Pseudo-Code:

for line in sys.stdin:

 doc = line.split('\t')

 doc_id = doc[0]

 features = tokenizeDoc(doc[1])

 D = len(features)

for feature in features:

```
print "W=", feature, "D_ID=", doc_id, "D=", D, "\t", "1"
```

```
print "W=*,D_ID=", doc_id, "\t", D    #Printing W* just to calculate avg  
#length in reducer, to avoid unnecessary parsing
```

```
print "D*=", "\t",
```

Output:

W=alpha,D_ID=id1,D=560 1

W=beta,D_ID=id1,D=560 1

W=zebra,D_ID=id1,D=560 1

W=*,D_ID=id1 560

D=* 1

W=alpha,D_ID=id2,D=100 1

W=beta,D_ID=id2,D=100 1

W=zabra,D_ID=id2,D=100 1

W=*,D_ID=id2 100

D=* 1

.

.

.

W=zebra,D_ID=idn,D=500 1

W=*,D_ID=idn 500

D=* 1

Sort After Mapper 1:

Output:

D=* 1

D=* 1

D=* 1

W=*,D_ID=id1 560

W=*,D_ID=id2 100

W=*,D_ID=idn 500

W=alpha,D_ID=id1,D=560 1

W=alpha,D_ID=id2,D=100 1

W=beta,D_ID=id1,D=560 1

W=beta,D_ID=id2,D=100 1

W=zebra,D_ID=id1,D=560 1

W=zabra,D_ID=id2,D=100 1

W=zebra,D_ID=idn,D=500 1

Reducer 1:

Pseudo-Code:

#Calculate avg doc length and reduce counts

total_doc_cnt = 0

total_word_cnt = 0

for line in sys.stdin:

 m = line.split("\t")

 event = m[0]

 cnt = int(m[1].replace("\n", ""))

 if "W=*" in event: #Using W* just to calculate avg document length

 total_doc_cnt += 1 #Not printing document length again,

 total_word_cnt += cnt #as it is already included in word string

 else:

 if event == event_prev:

 cnt_prev += cnt

 else:

 if event_prev != "":

 print event_prev, "\t", cnt_prev, "\n"

 cnt_prev = cnt

 event_prev = event

```

else:

    event_prev = event

    cnt_prev = cnt

#To print last event because it won't go in the if as there won't be any event
after that

print event_prev, "\t", cnt_prev, "\n"

#Print avg_length of the document

print "avgdl", "\t", (total_word_cnt/total_doc_cnt)

```

Output:

```

D=* 777

W=*,D_ID=id1 560

W=*,D_ID=id2 100

W=*,D_ID=idn 500

W=alpha,D_ID=id1,D=560 50

W=alpha,D_ID=id2,D=100 60

W=beta,D_ID=id1,D=560 70

W=beta,D_ID=id2,D=100 80

W=zebra,D_ID=id1,D=560 90

W=zabra,D_ID=id2,D=100 95

W=zebra,D_ID=idn,D=500 11

```

Mapper 2:

Pseudo-Code:

for line in sys.stdin:

```
    m = line.split("\t")
```

```
    event = m[0]
```

```
    cnt = int(m[1].replace("\n", ""))
```

```
    if "W=" in event and "Y=" in event:
```

```
        if event == event_prev:          #Print as it is and accumulate count
```

```
            print event, "\t", cnt, "\n"   #W=alpha,D_ID=id1,D=560    50
```

```
            cnt_prev += 1
```

```
    else:
```

```
        if event_prev != "":
```

```
            print event_prev, "\t", cnt_prev, "\n"
```

```
            word = extract_word_from_event(event_prev) #extract alpha from
W=alpha,D_ID=id1,D=560
```

```
            print "W=", word, "D_ID=*~t", cnt_prev #Print number of
documents containing word
```

```
            cnt_prev = 1
```

```
            event_prev = event
```

```

else:

    event_prev = event

    cnt_prev = 1

#To print last event because it won't go in the if as there won't be any event
after that

word = extract_word_from_event(event_prev) #extract alpha from
W=alpha,D_ID=id1,D=560

print "W=", word, "D_ID=*\t", cnt_prev #Print number of documents
containing word

```

Output:

```

D=* 777

W=*,D_ID=id1 560

W=*,D_ID=id2 100

W=*,D_ID=idn 500

W=alpha,D_ID=id1,D=560 50

W=alpha,D_ID=id2,D=100 60

W=alpha,D_ID=* 60

W=beta,D_ID=id1,D=560 70

W=beta,D_ID=id2,D=100 80

W=beta,D_ID=* 80

```


W=zebra,D_ID=id1,D=560 90

W=zabra,D_ID=id2,D=100 95

W=zebra,D_ID=idn,D=500 11

W=zebra,D_ID=*11

avgdl 999

Reducer 2:

Pseudo-Code:

#Sort will just bring n(wi) term at the beginning of every

W=word,D_ID=id,D=cnt

#so that it eases calculation in mapper 3 and we don't have to store n(wi) for every word

#W=alpha,D_ID=* will automatically come before

"W=alpha,D_ID=id1,D=560" because of sort

Output:

avgdl 999

D=* 777

W=alpha,D_ID=* 60

W=alpha,D_ID=id1,D=560 50

W=alpha,D_ID=id2,D=100 60

W=beta,D_ID=* 80

W=beta,D_ID=id1,D=560 70

W=beta,D_ID=id2,D=100 80

W=zebra,D_ID=*11

W=zebra,D_ID=id1,D=560 90

W=zabra,D_ID=id2,D=100 95

W=zebra,D_ID=idn,D=500 11

Mapper 3:

Pseudo-Code:

#This will store query in memory and stream throught the conunts stored from

#reducer 2 and calculte score for every di and qi

#Assume counts from mapper 2 are coming from standard input

Q = file.read("Query")

for line in sys.stdin:

 m = line.split("\t")

 event = m[0]

 cnt = int(m[1].replace("\n", ""))

 if event == "avgdl":

 avgdl = cnt #it will come first because of sorting

 if "D_ID=*" in line:

```
n_wi = cnt
```

```
else:
```

```
word, did, D = extract_word_id_D_from_event(event) #Parsing based on  
#commas and equal to sign
```

```
if word in Q:
```

```
IDF = log((N - n(qi)+0.5) / (n(qi) + 0.5))
```

```
score = (IDF * cnt * (k1 + 1))/(cnt + k1*(1-b+(b*D/avgdl)))
```

```
print did,word, "\t", score    #d1,q2 0.08
```

```
                                #d2,q2 0.04
```

Output:

```
d1, q2    0.08
```

```
d2, q2    0.04
```

```
d3, q3    0.03
```

```
d4, q2    0.02
```

```
d1, q1    0.08
```

```
d2, q1    0.04
```

```
d3, q1    0.03
```

```
d4, q1    0.02
```

Reducer 3:

Pseudo-Code:

```
#As every statement in mapper output start with document id,

# The input to reducer will be sorted based on document id, hence score for

# all queries for document will come sequentially

#Reducer will just add the scores and print final score for each document

d_prev = ""

for line in sys.stdin:

    m = line.split("\t")

    event = m[0]

    d,q = extract_dq_from_event(event)

    score = int(m[1].replace("\n",""))

    if d == d_prev:

        score_prev += score

    else:

        if d_prev != "":

            print d, "score(", d, ",Q") =, score_prev

        cnt_prev = score
```

```
d_prev = d
```

```
else:
```

```
d_prev = d
```

```
score_prev = score
```

```
#To print last event because it won't go in the if as there won't be any event  
after that
```

```
print d, "score(", d, ",Q") =, score_prev
```

Output:

```
d1, score(d1, Q) = 0.08
```

```
d2, score(d2, Q) = 0.09
```

```
d3, score(d1, Q) = 0.08
```

```
d4, score(d2, Q) = 0.09
```